



# Integrating LLVM with the PGI Fortran, C and C++ compilers

Third Workshop on the LLVM Compiler Infrastructure in HPC

Doug Miles, PGI Compilers & Tools, NVIDIA Corporation, 14 November, 2016

# PGI Fortran , C & C++ Compilers

Optimizing, SIMD Vectorizing, OpenMP

## Accelerated Computing Features

OpenACC Directives

CUDA Fortran

## Multi-Platform Solution

x86-64 and OpenPOWER CPUs,  
Tesla and Radeon GPUs

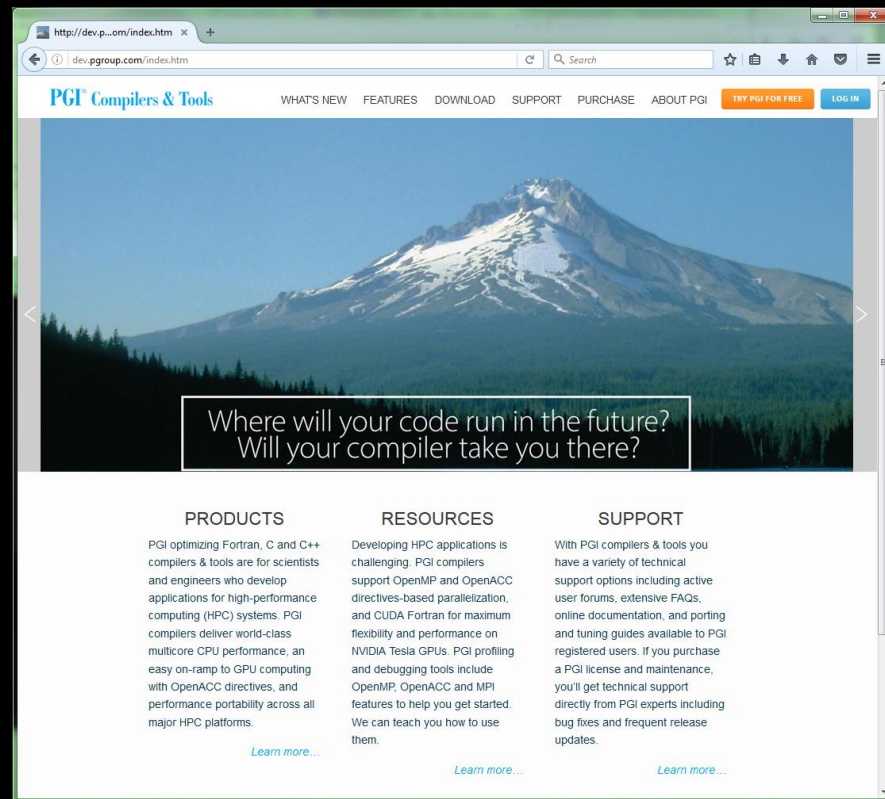
Supported on Linux, macOS, Windows

## MPI/OpenMP/OpenACC Tools

PGDBG® debugger

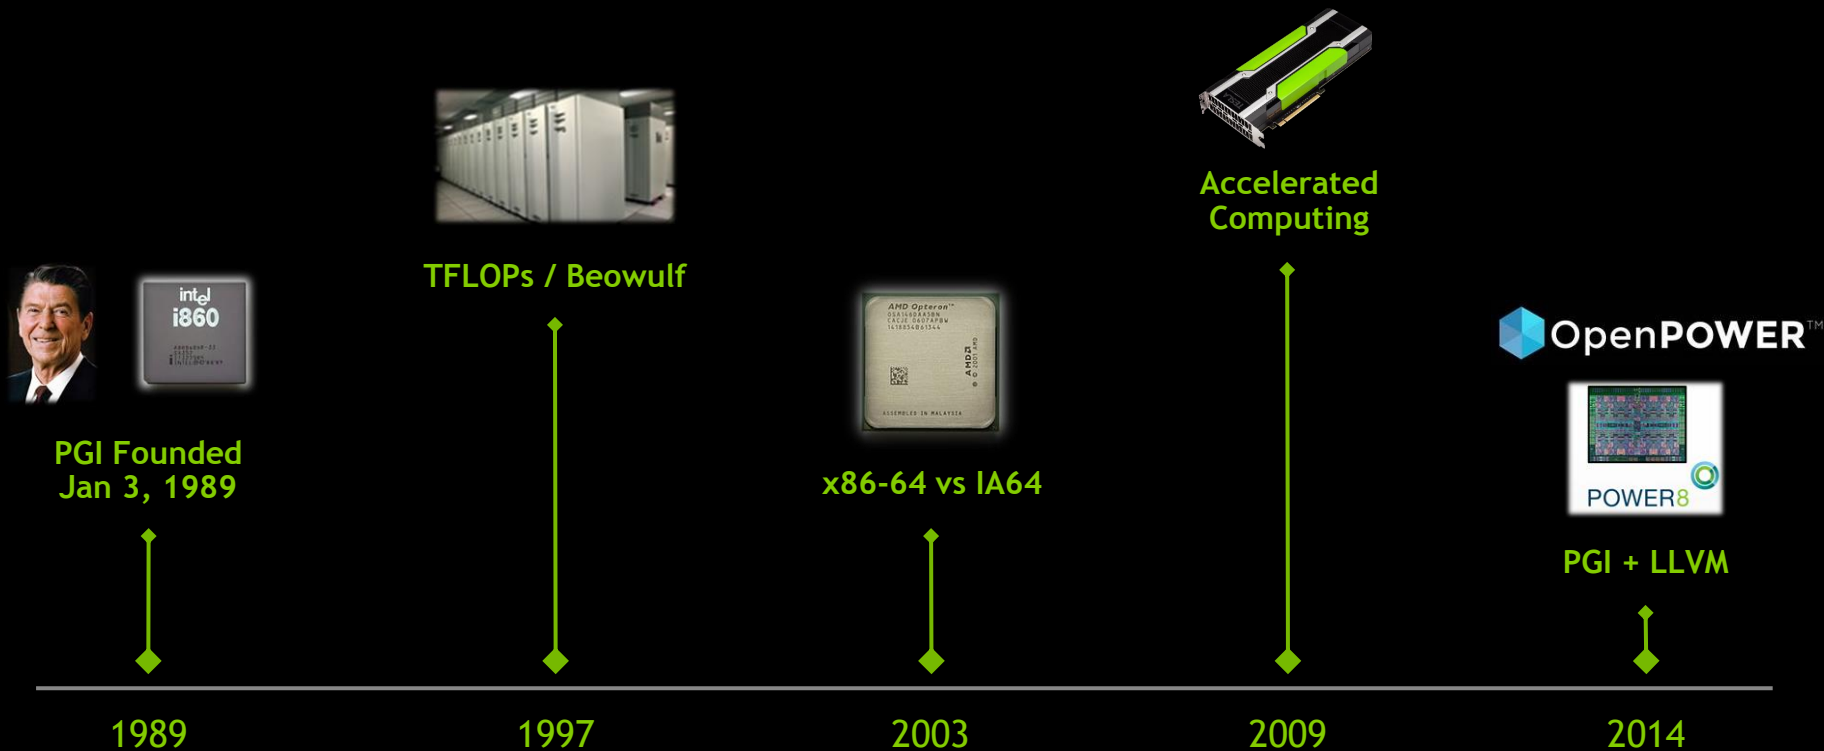
PGPROF® profiler

Interoperable with DDT, Totalview



[www.pgroup.com](http://www.pgroup.com)

# Riding Waves of Disruption

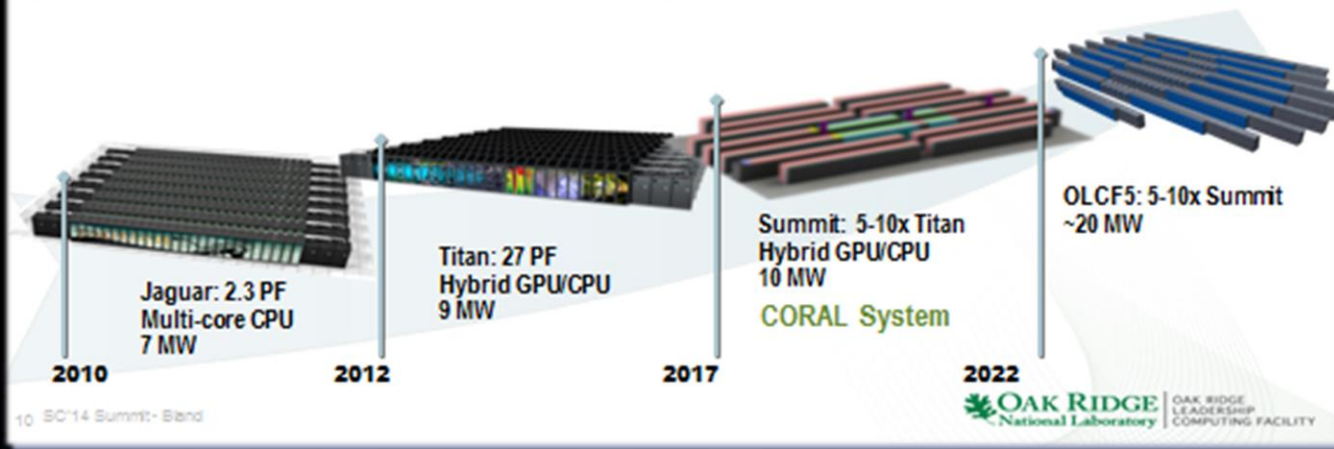


# ORNL Leadership Computing Systems, CORAL

**Our Science requires that we continue to advance OLCF's computational capability over the next decade on the roadmap to Exascale.**

Since clock-rate scaling ended in 2003, HPC performance has been achieved through increased parallelism. Jaguar scaled to 300,000 cores.

Titan and beyond deliver hierarchical parallelism with very powerful nodes. MPI plus thread level parallelism through OpenACC or OpenMP plus vectors



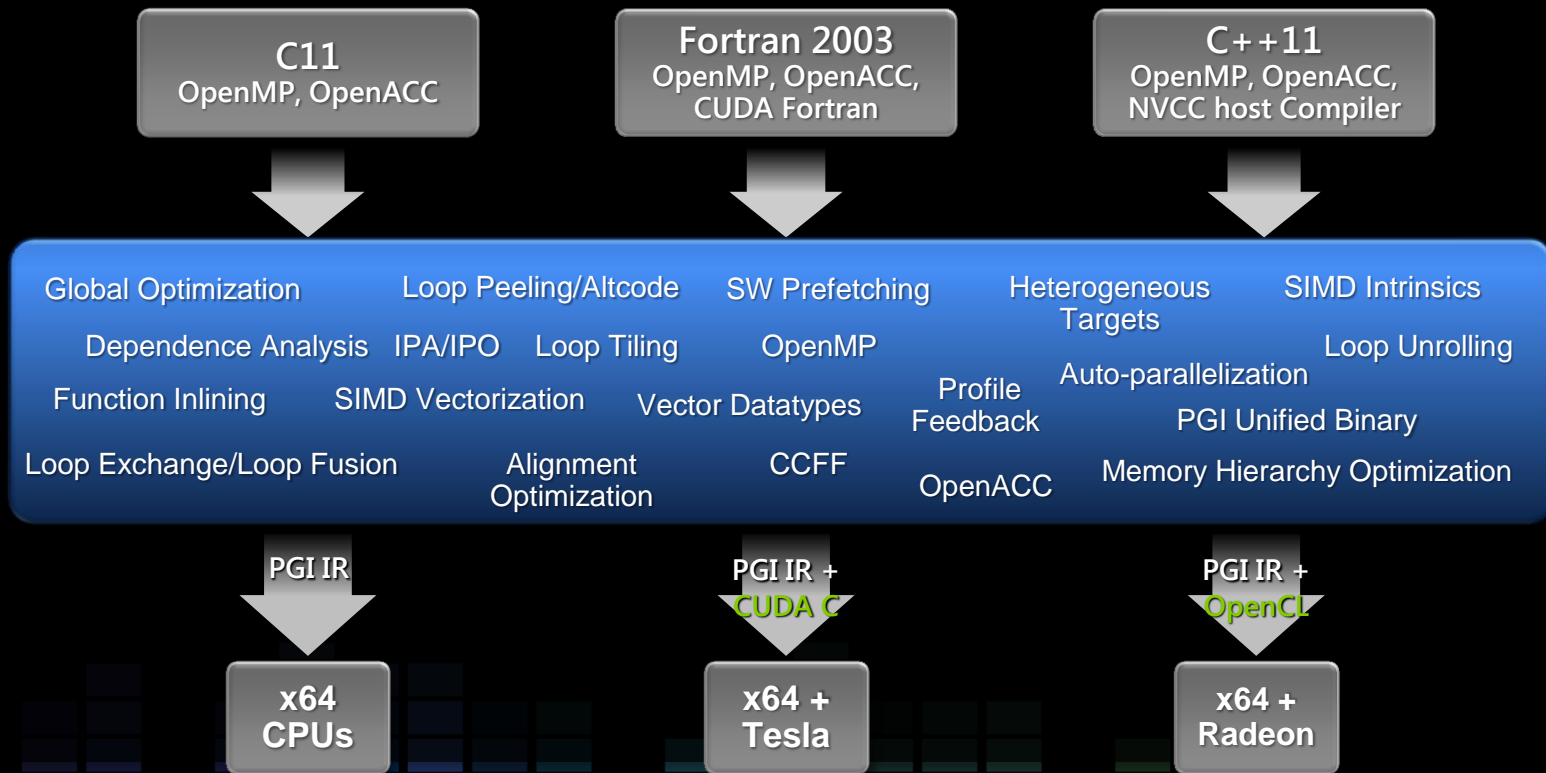
# PGI and CORAL



"Porting and optimizing production HPC applications from one platform to another can be one of the most significant costs in the adoption of breakthrough hardware technologies. The PGI compiler has been our primary compiler on Jaguar and Titan since 2005. Having the PGI compiler suite available in the POWER environment will provide continuity and facilitate code portability of existing CPU-only and GPU-enabled Titan applications to our next major system. "

— *Buddy Bland, Titan Project Director, Oak Ridge National Lab*

# PGI Compilers 2014 ...



# LLVM: Community Power

Contributing Organizations

Processors

2000



10

15



2005



21

Google



2010



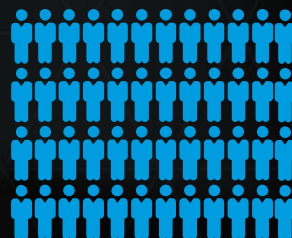
178

SAMSUNG Microsoft

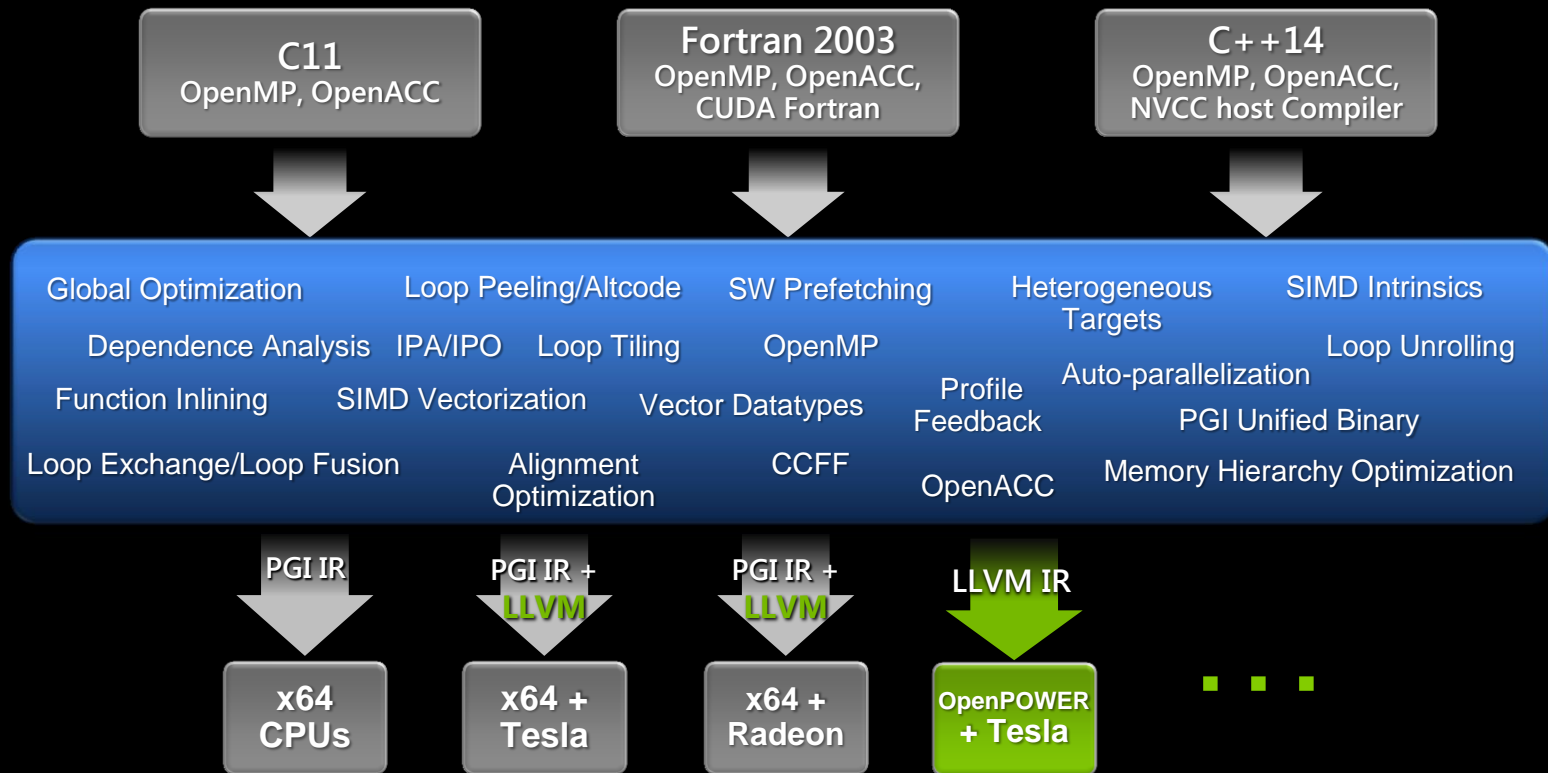


475

Active LLVM Contributors



# PGI Compilers 2016 ...





# PGI for OpenPOWER+Tesla

Fortran 2003, C11, C++14 compilers,  
PGPROF profiler

CUDA Fortran, OpenACC, OpenMP,  
NVCC host compiler

Integrated with LLVM for  
OpenPOWER code generation

First production release now available



PGI®

x86

Recompile



PGI®

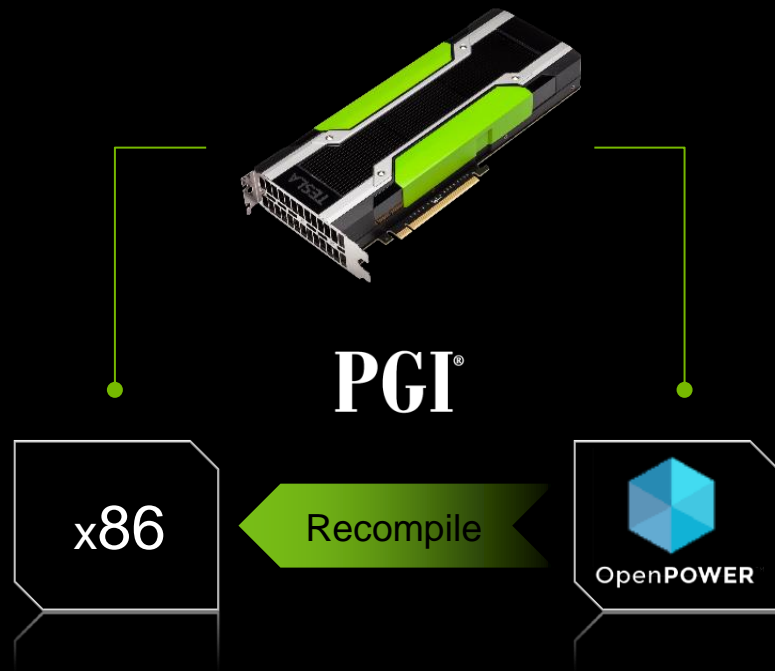
# PGI for OpenPOWER+Tesla

Fortran 2003, C11, C++14 compilers,  
PGPROF profiler

CUDA Fortran, OpenACC, OpenMP,  
NVCC host compiler

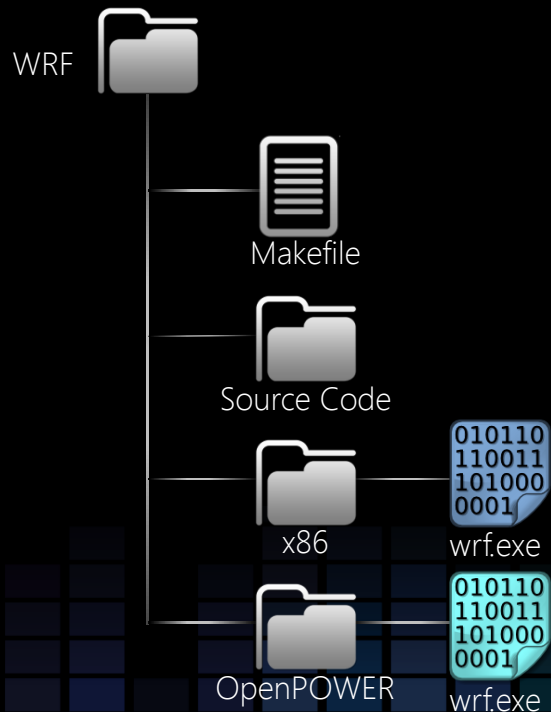
Integrated with LLVM for  
OpenPOWER code generation

First production release now available



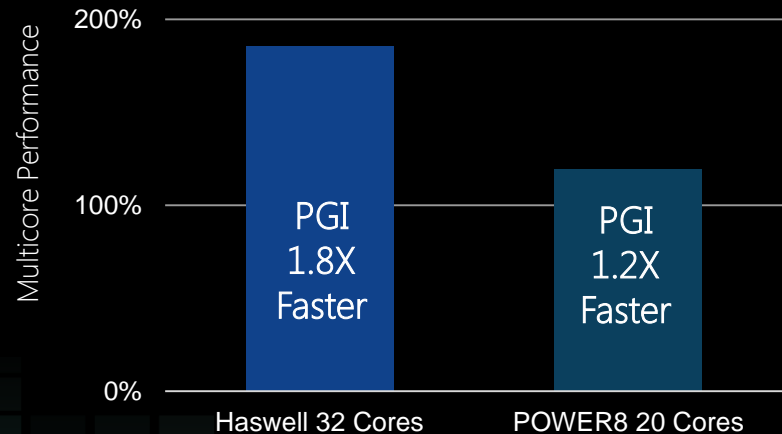
# Porting an 800K line HPC Application from x86 to OpenPOWER

Recompile ...



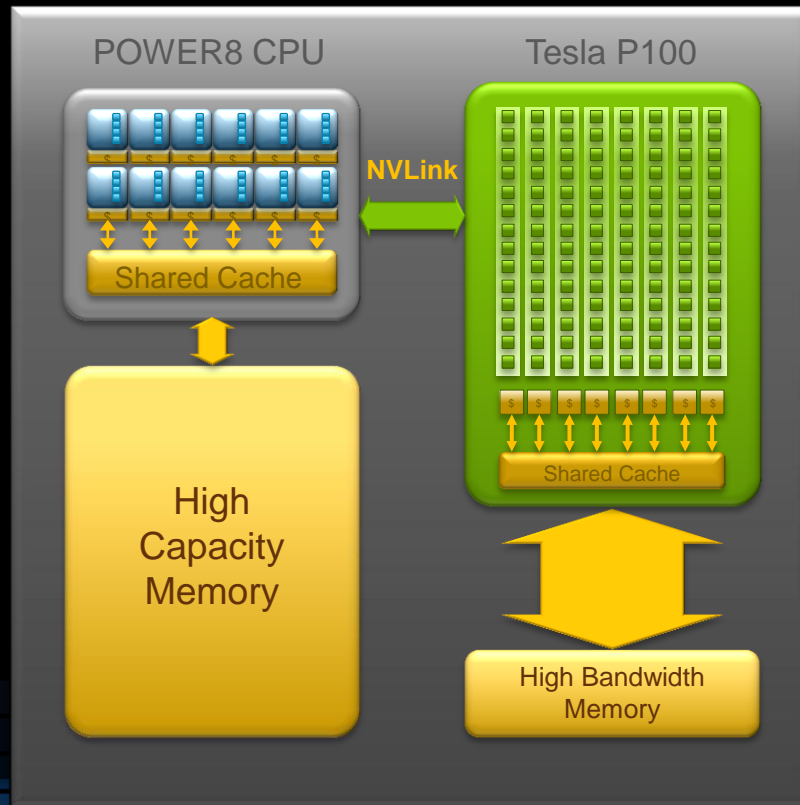
Run ...

WRF 3.8.1 OpenMP Performance  
PGI 16.10 vs GNU 6.1

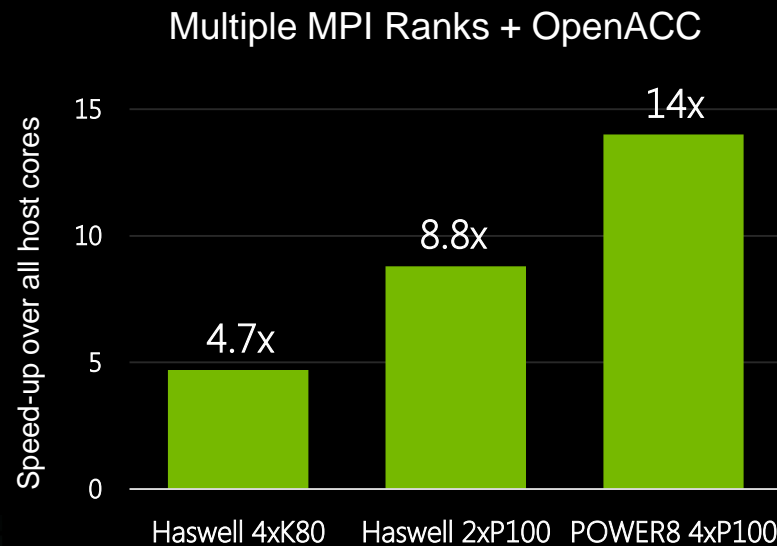
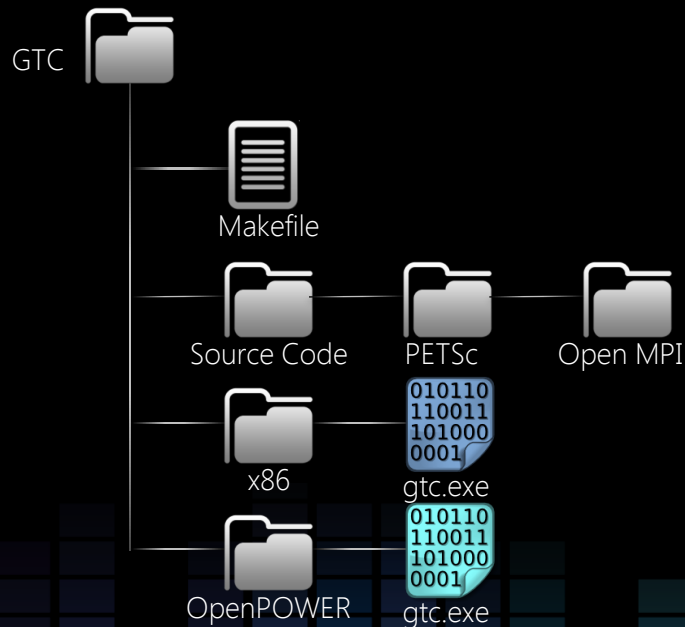


x86 CPU: Intel Xeon E5-2698 v3, 2 sockets, 32 cores  
OpenPOWER CPU: IBM 8247-42L POWER8E, 4 sockets, 20 cores  
PGI options: -fast -Mstack\_arrays -mp  
GNU options: -O3 -funroll-loops -fpeel-loops -fopenmp

# OpenPOWER+Tesla HPC Node



# Porting the Gyrokinetic Toroidal Code (GTC) from Xeon+Tesla to OpenPOWER+Tesla using OpenACC



X86 CPU: Intel Xeon E5-2698 v3,  
POWER CPU: IBM POWER8NVL

# OpenACC Directives

Manage  
Data  
Movement

```
#pragma acc data copyin(a,b) copyout(c)
{
```

Initiate  
Parallel  
Execution

```
...
#pragma acc parallel
{
  #pragma acc loop gang vector
  for (i = 0; i < n; ++i) {
    c[i] = a[i] + b[i];
    ...
  }
}
```

Optimize  
Loop  
Mappings

```
...
}
```

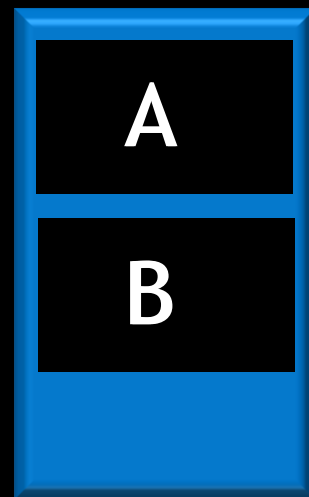
**OpenACC**  
Directives For Accelerators

- Incremental
- Single source
- Interoperable
- Performance portable
- CPU, GPU, Manycore

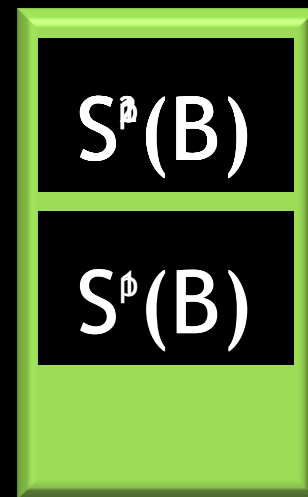


# OpenACC in a Nutshell

```
➡ ...  
➡ #pragma acc data copy(b[0:n][0:m]) \  
    create(a[0:n][0:m])  
  
{  
➡ for (iter = 1; iter <= p; ++iter){  
➡   #pragma acc kernels  
   {  
➡     for (i = 1; i < n-1; ++i){  
       for (j = 1; j < m-1; ++j){  
         a[i][j]=w0*b[i][j]+  
➡           w1*(b[i-1][j]+b[i+1][j]+  
               b[i][j-1]+b[i][j+1])+  
           w2*(b[i-1][j-1]+b[i-1][j+1]+  
               b[i+1][j-1]+b[i+1][j+1]);  
       } }  
➡     for( i = 1; i < n-1; ++i )  
➡       for( j = 1; j < m-1; ++j )  
➡         b[i][j] = a[i][j];  
➡     }  
➡   }  
➡   ...  
}
```



Host  
Memory



Accelerator  
Memory

# OpenACC for Multicore CPUs & GPUs

```
!$acc kernels loop
do j = 1, m
  do i = 1, n
    a(j,i) = b(j,i)*alpha + c(i,j)*beta
  enddo
enddo
```

CPU

GPU

```
% pgfortran a.f90 -ta=multicore -c -Minfo
sub:
```

```
10, Loop is parallelizable
    Generating Multicore code
10, !$acc loop gang
11, Loop is parallelizable
```

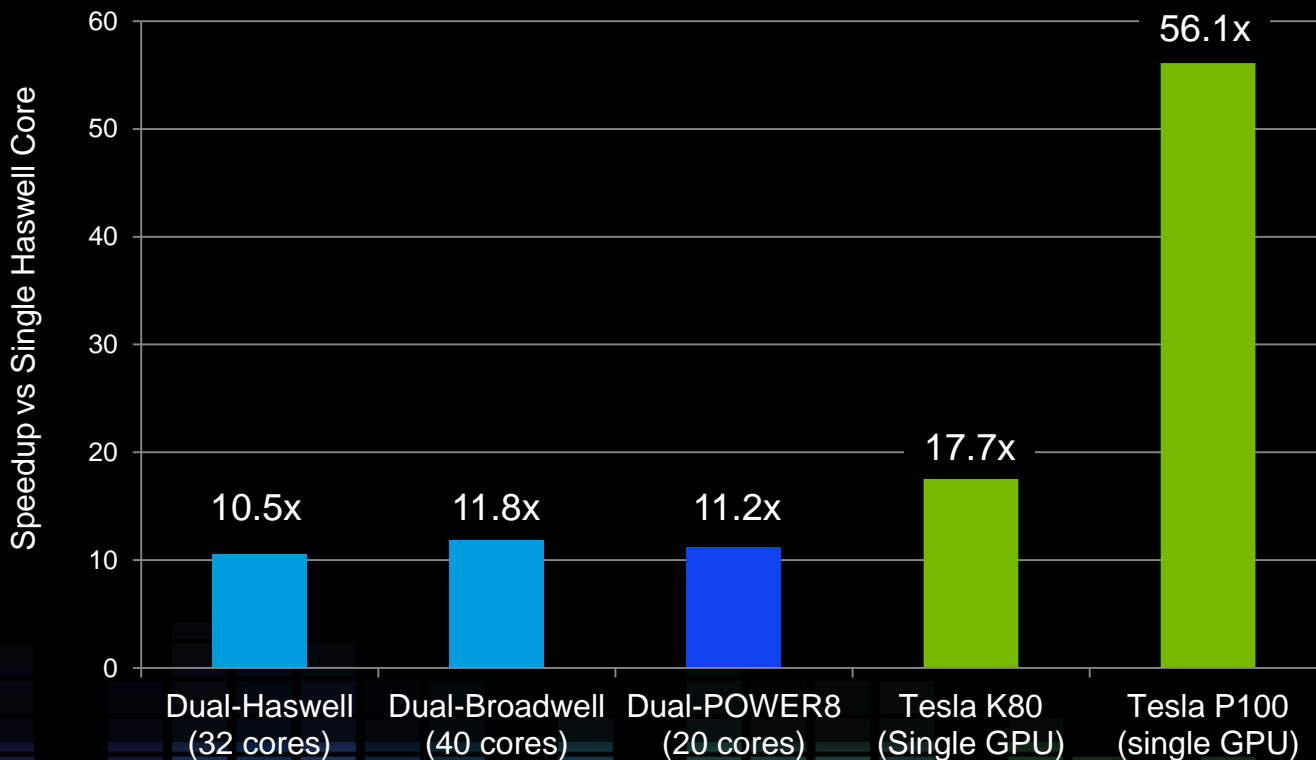
```
% pgfortran a.f90 -ta=tesla -c -Minfo
sub:
```

```
10, Loop is parallelizable
11, Loop is parallelizable
    Accelerator kernel generated
    Generating Tesla code
10, !$acc loop gang, vector(4)
11, !$acc loop gang, vector(32)
```



# PGI OpenACC - SPEC ACCEL 1.0 Benchmarks

Geometric mean across all 15 benchmarks



# CUDA Fortran for Tesla GPUs

```
real, device, allocatable, dimension(:,:) ::  
    Adev,Bdev,Cdev
```

```
...
```

```
→ allocate (Adev(N,M), Bdev(M,L), Cdev(N,L))
```

```
→ Adev = A(1:N,1:M)
```

```
→ Bdev = B(1:M,1:L)
```

```
→ call mm_kernel <<<dim3(N/16,M/16),dim3(16,16)>>>  
    ( Adev, Bdev, Cdev, N, M, L )
```

```
→ C(1:N,1:L) = Cdev
```

```
→ deallocate ( Adev, Bdev, Cdev )
```

```
→ ...
```

CPU Code

PGI®

```
attributes(global) subroutine mm_kernel  
    ( A, B, C, N, M, L )  
    real :: A(N,M), B(M,L), C(N,L), Cij  
    integer, value :: N, M, L  
    integer :: i, j, kb, k, tx, ty  
    real, shared :: Asub(16,16),Bsub(16,16)  
    tx = threadidx%x  
    ty = threadidx%y  
    i = blockidx%x * 16 + tx  
    j = blockidx%y * 16 + ty  
    Cij = 0.0  
    do kb = 1, M, 16  
        Asub(tx,ty) = A(i,kb+tx-1)  
        Bsub(tx,ty) = B(kb+ty-1,j)  
        call syncthreads()  
        do k = 1,16  
            Cij = Cij + Asub(tx,k) * Bsub(k,ty)  
        enddo  
        call syncthreads()  
    enddo  
    C(i,j) = Cij  
end subroutine mmul_kernel
```

Tesla Code

# !\$CUF kernel directives

```

module madd_device_module
  use cudafor
contains
  subroutine madd_dev(a,b,c,sum,n1,n2)
    real,dimension(:,:),device :: a,b,c
    real :: sum
    integer :: n1,n2
    type(dim3) :: grid, block
    !$cuf kernel do (2) <<<(*,*),(32,4)>>>
  {
    do j = 1,n2
      do i = 1,n1
        a(i,j) = b(i,j) + c(i,j)
        sum = sum + a(i,j)
      enddo
    enddo
  }
end subroutine
end module

```

Equivalent  
hand-written  
CUDA kernels



```

module madd_device_module
  use cudafor
  implicit none
contains
  attributes(global) subroutine madd_kernel(a,b,c,blocksum,n1,n2)
    real, dimension(:,:) :: a,b,c
    real, dimension(:) :: blocksum
    integer, value :: n1,n2
    integer :: i,j,tindex,tneighbor,bindex
    real :: mysum
    real, shared :: bsum(256)
    ! Do this thread's work
    mysum = 0.0
    do j = threadidx%y + (blockidx%y-1)*blockdim%y, n2, blockdim%y*griddim%y
      do i = threadidx%x + (blockidx%x-1)*blockdim%x, n1, blockdim%x*griddim%x
        a(i,j) = b(i,j) + c(i,j)
        mysum = mysum + a(i,j) ! accumulates partial sum per thread
      enddo
    enddo
    ! Now add up all partial sums for the whole thread block
    ! Compute this thread's linear index in the thread block
    ! We assume 256 threads in the thread block
    tindex = threadidx%x + (threadidx%y-1)*blockdim%x
    ! Store this thread's partial sum in the shared memory block
    bsum(tindex) = mysum
    call syncthreads()
    ! Accumulate all the partial sums for this thread block to a single value
    tneighbor = 128
    do while( tneighbor >= 1 )
      if( tindex <= tneighbor ) &
        bsum(tindex) = bsum(tindex) + bsum(tindex+tnneighbor)
      tneighbor = tneighbor / 2
      call syncthreads()
    enddo
    ! Store the partial sum for the thread block
    bindex = blockidx%x + (blockidx%y-1)*griddim%x
    if( tindex == 1 ) blocksum(bindex) = bsum(1)
  end subroutine

  ! Add up partial sums for all thread blocks to a single cumulative sum
  attributes(global) subroutine madd_sum_kernel(blocksum,dsum,nb)
    real, dimension(:) :: blocksum
    real :: dsum
    integer, value :: nb
    real, shared :: bsum(256)
    integer :: tindex,tnneighbor,i
    ! Again, we assume 256 threads in the thread block
    ! accumulate a partial sum for each thread
    tindex = threadidx%x
    bsum(tindex) = 0.0
    do i = tindex, nb, blockdim%x
      bsum(tindex) = bsum(tindex) + blocksum(i)
    enddo
    call syncthreads()
    ! This code is copied from the previous kernel
    ! Accumulate all the partial sums for this thread block to a single value
    ! Since there is only one thread block, this single value is the final result
    tneighbor = 128
    do while( tneighbor >= 1 )
      if( tindex <= tneighbor ) &
        bsum(tindex) = bsum(tindex) + bsum(tindex+tnneighbor)
      tneighbor = tneighbor / 2
      call syncthreads()
    enddo
    if( tindex == 1 ) dsum = bsum(1)
  end subroutine

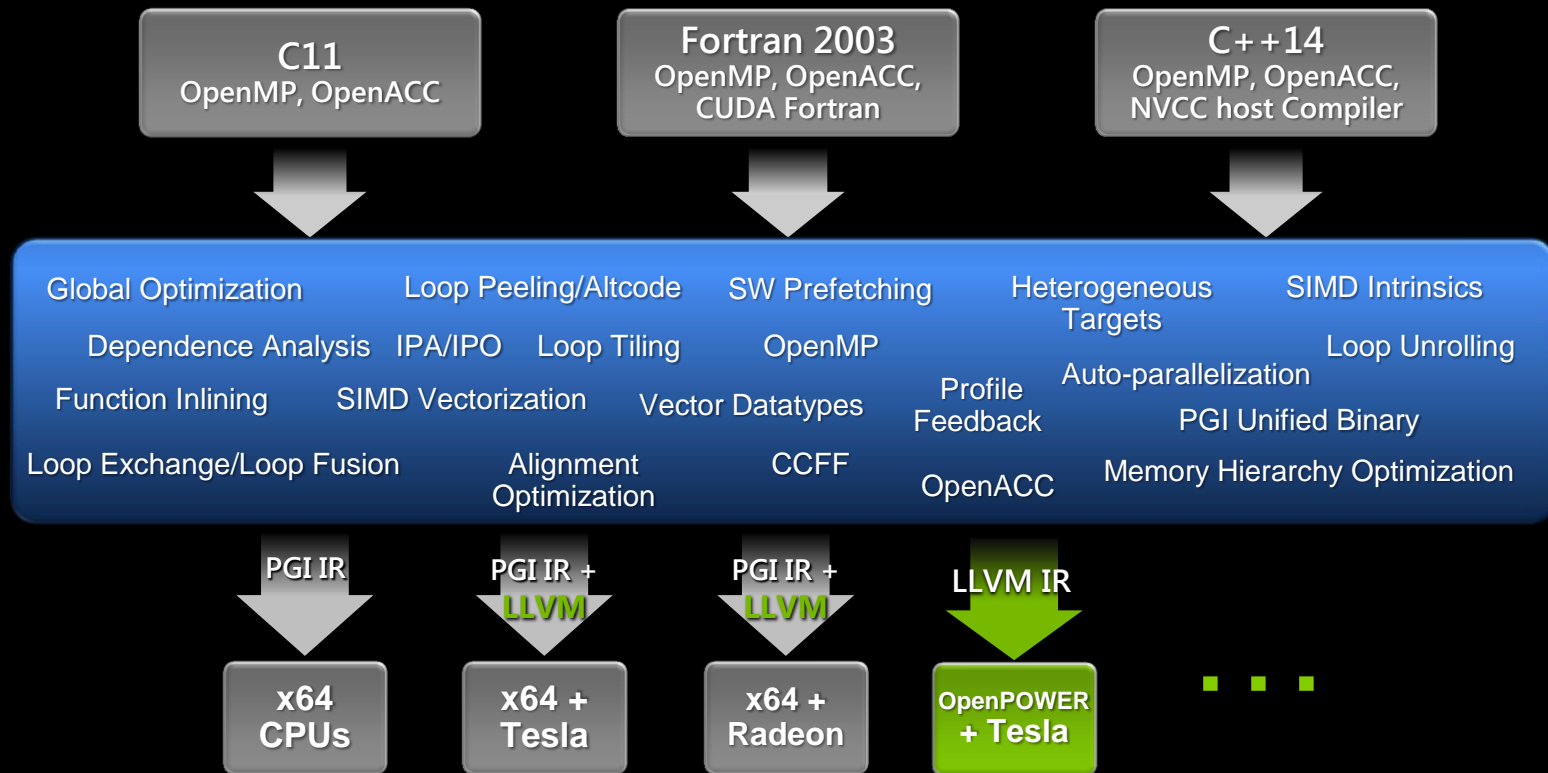
  subroutine madd_dev(a,b,c,dsum,n1,n2)
    real, dimension(:,:), device :: a,b,c
    real, device :: dsum
    real, dimension(:), allocatable, device :: blocksum
    integer :: n1,n2,nb
    type(dim3) :: grid, block
    integer :: r
    ! Compute grid/block size; block size must be 256 threads
    grid = dim3((n1+31)/32, (n2+7)/8, 1)
    block = dim3(32,8,1)
    nb = grid%x * grid%y
    allocate(blocksum(1:nb))
    call madd_kernel<< grid, block >>>(a,b,c,blocksum,n1,n2)
    call madd_sum_kernel<< 1, 256 >>>(blocksum,dsum,nb)
    r = cudaThreadSynchronize() ! don't deallocate too early
    deallocate(blocksum)
  end subroutine
end module

```

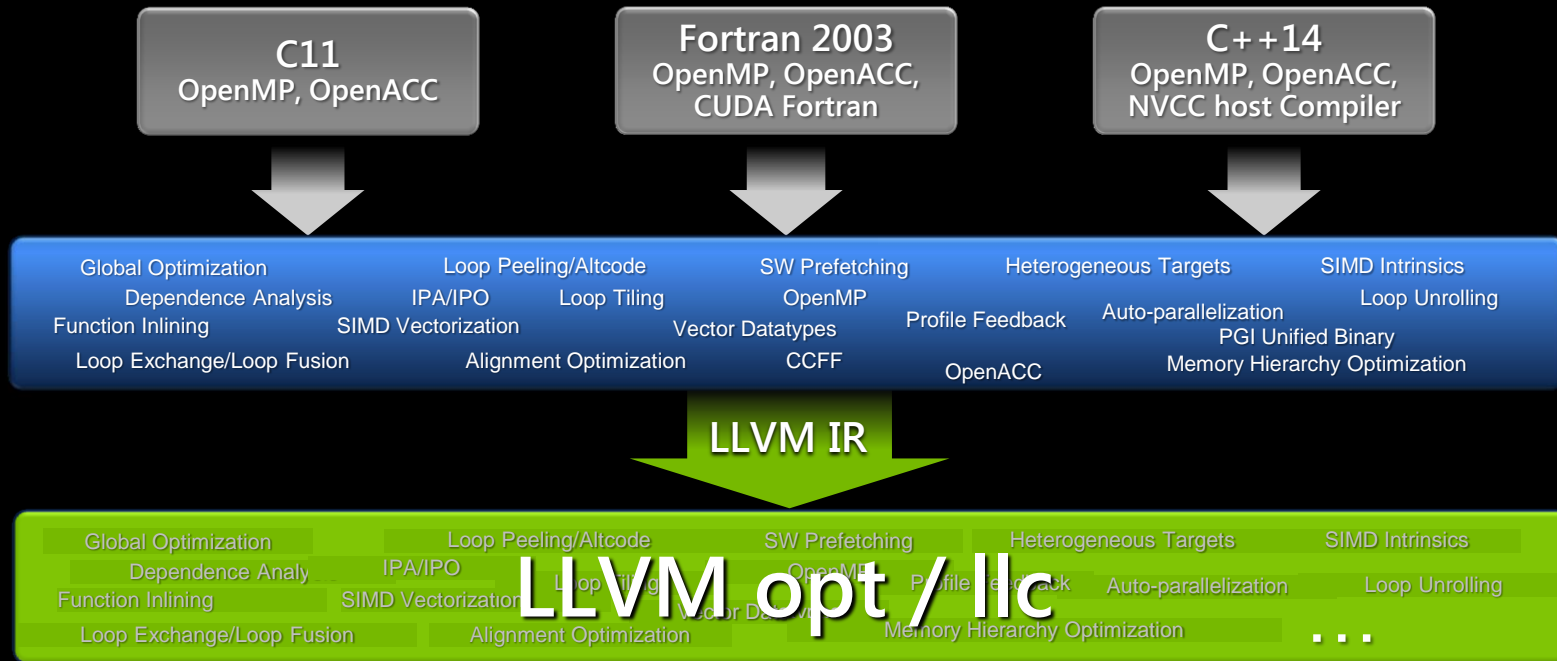
# PGI + LLVM Integration



# PGI Compilers 2016 ...



# LLVM is not a code generator ...



it is "... a collection of modular and reusable compiler and toolchain technologies."

# Integrating LLVM into the PGI Compilers

- PGI ILI -> LLVM IR bridge, CPU-side and GPU-side
- C/C++/Fortran language support, scalar code generation
- Target independent vectorizer
- OpenMP re-implementation, SMP auto-parallelization
- Enabling OpenACC and CUDA Fortran
- Integration, Testing, Documentation
- Dovetailing PGI optimizer and LLVM opt



# Target-independent vectorizer

```
subroutine add_exp(n,a,b,c)
  integer n
  real*8, dimension(n) :: a,b,c
  integer i

  do i = 1,n
    a(i) = b(i) + exp(c(i))
  enddo
end subroutine
```

L.LB1\_428:

```
. . .
%14 = getelementptr i8, i8* %11, i64 %13      # compute c(i) location
%15 = bitcast i8* %14 to <4 x double>*      # cast c(i) as vector of 4 doubles
%16 = load <4 x double>, <4 x double>* %15, align 8
%17 = bitcast <4 x double> (...) * @__gvd_exp4 to <4 x double> (<4 x double>)*
%18 = call <4 x double> %17 (<4 x double> %16) # call vector exp on 4 doubles
. . .
%22 = getelementptr i8, i8* %19, i64 %21      # compute b(i) location
%23 = bitcast i8* %22 to <4 x double>*      # cast b(i) as vector of 4 doubles
%24 = load <4 x double>, <4 x double>* %23, align 8
%25 = fadd <4 x double> %18, %24             # add b(i) to exp(c(i))
. . .
%28 = getelementptr i8, i8* %26, i64 %27      # compute a(i) location
%29 = bitcast i8* %28 to <4 x double>*      # cast a(i) as vector of 4 doubles
store <4 x double> %25, <4 x double>* %29, align 1 # store result to a(i)
. . .
br i1 %33, label %L.LB1_428, label %L.LB1_473 # loop
...
declare <4 x double> @__gvd_exp4(...)
```

Leveraging LLVM Vector Data Types



# Target-independent vectorizer

```
subroutine add_exp(n,a,b,c)
  integer n
  real*8, dimension(n) :: a,b,c
  integer i

  do i = 1,n
    a(i) = b(i) + exp(c(i))
  enddo

end subroutine
```

```
        leal    -3(%rbx), %r12d
.LBB0_3:
        vmovupd (%r14,%rbp), %ymm0
        callq   __gvd_exp4
        vaddpd  (%r15,%rbp), %ymm0, %ymm0
        vmovupd %ymm0, (%r13,%rbp)
        addq    $32, %rbp
        addl    $-4, %r12d
        testl   %r12d, %r12d
        jg      .LBB0_3
```

**x86-64 AVX-256**

```
.LBB0_4:
        lxvd2x  0, 30, 24
        stxvd2x 0, 1, 23
        ori     2, 2, 0
        lxvd2x  0, 1, 23
        ld      4, 40(1)
        ld      3, 32(1)
        xxswapd 34, 0
        bl      __gvd_exp2
        nop
        lxvd2x  0, 29, 24
        addi    22, 22, -2
        cmpwi   22, 0
        xxswapd 0, 0
        xvaddpd 0, 34, 0
        xxswapd 0, 0
        stxvd2x 0, 28, 24
        addi    24, 24, 16
        bgt     0, .LBB0_4
```

**OpenPOWER VSX**

# Outlining parallel regions

```
subroutine add_exp(n,a,b,c)
  integer n
  real*8, dimension(n) :: a,b,c
  integer i
!$omp parallel do
  do i = 1,n
    a(i) = b(i) + exp(c(i))
  enddo
end subroutine
```

```
## lineno: 5
..LN1:
    movq    .STATICS1(%rip), %rdi
    movl    $2, %esi
    vzeroupper
    .p2align    4,,1
    call    _mp_penter
    vzeroupper
    .p2align    4,,1
    call    _mp_lcpu
    movl    %eax, 268(%rsp)
    vzeroupper
    .p2align    4,,1
    call    _mp_ncpus
    movl    268(%rsp), %ecx
    . . .
    # Execute SIMD vector loop
    # in parallel
    . . .
    call    _mp_pexit
    movq    -72(%rbp), %r15
    movq    -64(%rbp), %r14
```

```
. . .
xorl    %edi, %edi
callq   __kmpc_global_thread_num
movl    (%rbx), %eax
movl    %eax, 4(%rsp)
movq    %rbx, 16(%rsp)
movq    %r12, 24(%rsp)
leaq    4(%rsp), %rax
movq    %rax, 32(%rsp)
movq    %r15, 40(%rsp)
movq    %r14, 48(%rsp)
leaq    8(%rsp), %rcx
movl    $0, %edi
movl    $1, %esi
movl    $add_exp__1F1L5_, %edx
xorl    %eax, %eax
callq   __kmpc_fork_call
addq    $56, %rsp
popq    %rbx
popq    %r12
popq    %r14
popq    %r15
retq
```

# PGI + LLVM to do list

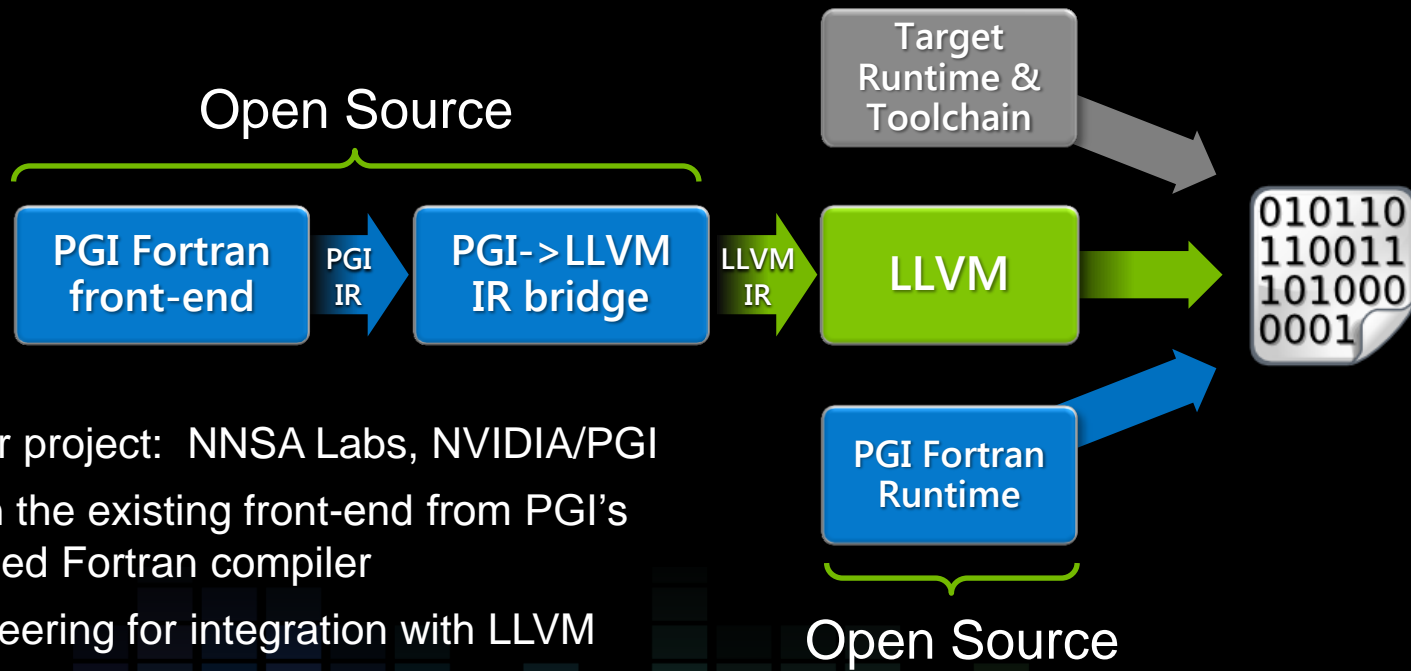
- Fortran DWARF generation
- OpenPOWER performance analysis
- OpenMP performance tuning, OpenMP 4.5
- PGI vectorizer performance tuning
- Dovetailing PGI optimizer and LLVM opt
- POWER9 128-bit IEEE floating-point support



Flang



# An open source Fortran front-end for LLVM a.k.a. the Flang project



- Multi-year project: NNSA Labs, NVIDIA/PGI
- Based on the existing front-end from PGI's widely-used Fortran compiler
- Re-engineering for integration with LLVM
- Develop CLANG-quality Fortran msg facility

# Many Stakeholders, Many Goals

LANL

New developer productive in source base in 4 – 8 weeks

Sandia

Single-thread/SIMD and OpenMP 3.1 performance

LLNL

OpenMP 4.x features, GPU and OpenPOWER support

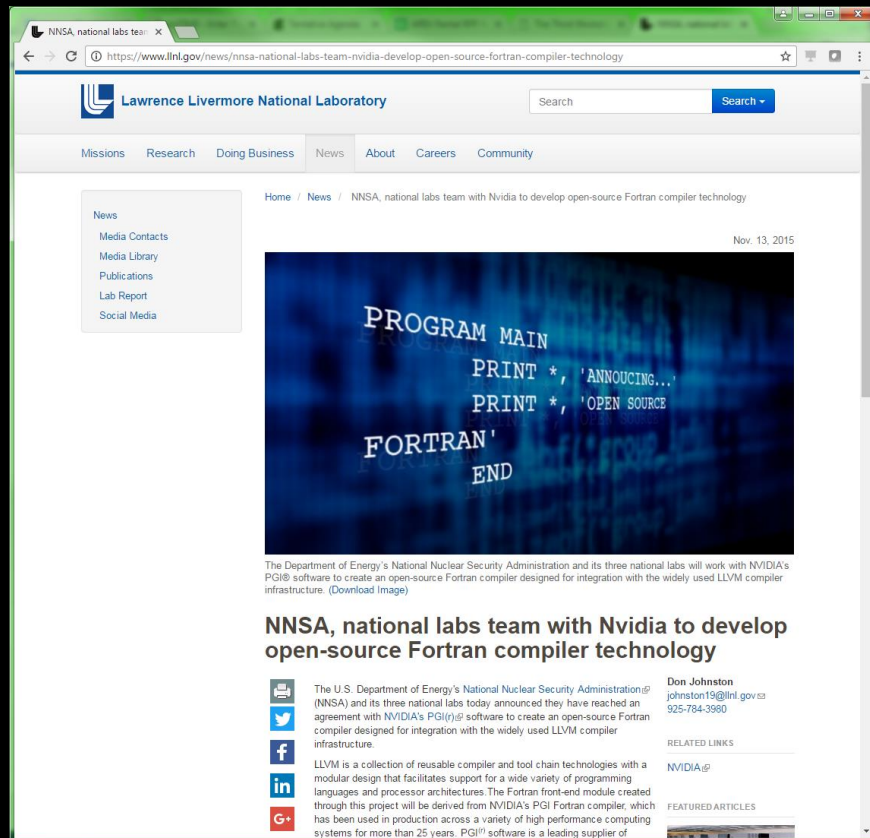
NVIDIA

Accelerate Fortran features support, PGI interoperability

Everyone

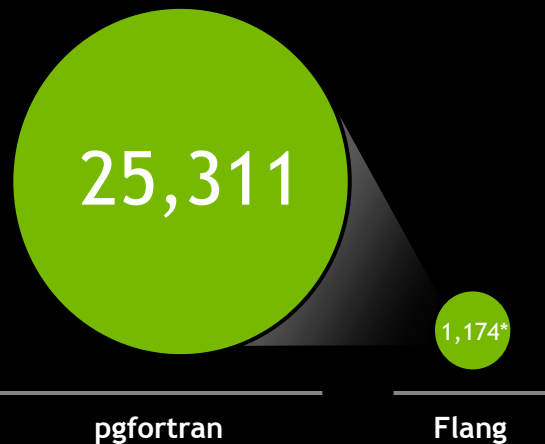
Adoption by both the HPC and LLVM communities

ANL, IBM, ARM Ltd, ORNL, Codethink, ...



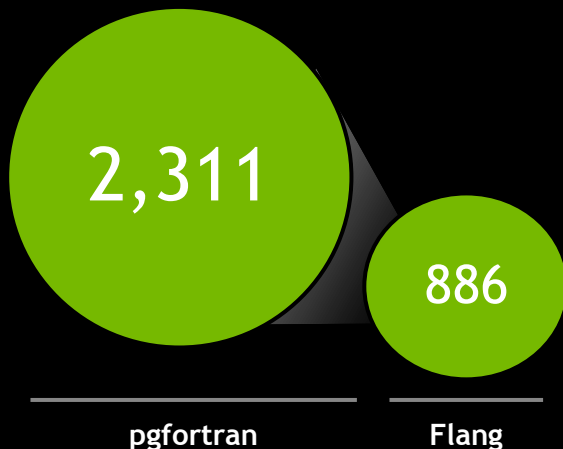
# Creating the initial Flang source base

- Front-end 85%
- Runtime Libraries 15%



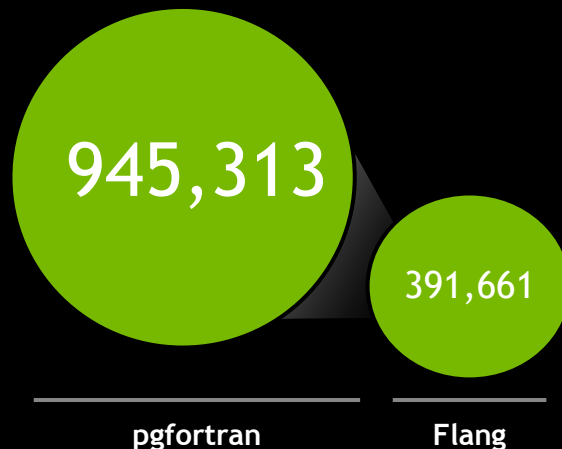
95% fewer #ifdefs

- Front-end 85%
- Runtime Libraries 15%



62% fewer files

- Front-end 80%
- Runtime Libraries 20%



59% fewer LOC

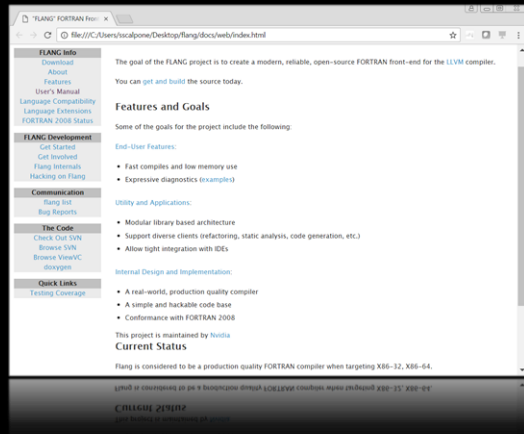
\*Clang has 212 #ifdefs in lib, include, tools

# Flang Development Status

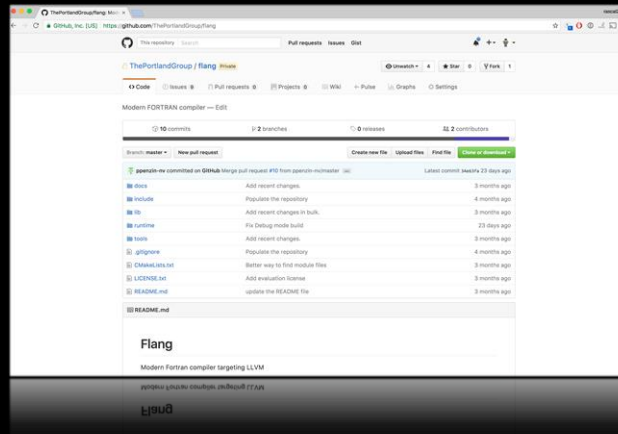
- Source code clean-up, refactoring & documentation ongoing
- Vendor neutrality nearly complete
- Frequent source and Flang binary updates to partners
- Passes most PGI Fortran Linux/x86 QA tests
- SIMD vectorization via the LLVM vectorizer, tuning ongoing
- Most of OpenMP 4.5 is implemented (CPU-side only)



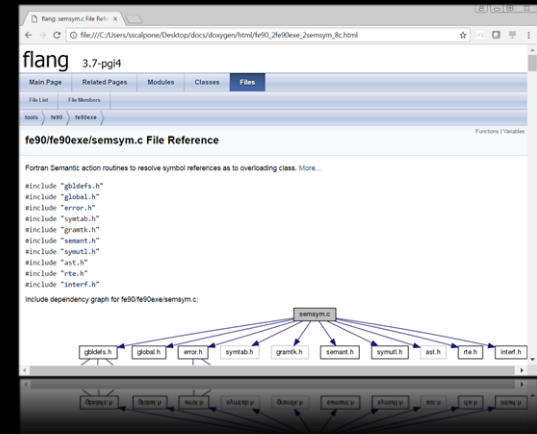
# Flang Source Code



Home page



Github



Doxygen

# Flang Single-core Performance

SPEC CPU 2006 Fortran codes, all times in seconds, 1 Haswell core

|               | PGI FORTRAN<br>16.10 | GFORTTRAN<br>6.1 | FLANG DEV<br>LLVM 3.9 |
|---------------|----------------------|------------------|-----------------------|
| 410.bwaves    | 182s                 | 220s             | 251s                  |
| 416.gamess    | 507s                 | Fails            | 475s                  |
| 434.zeusmp    | 183s                 | 221s             | 240s                  |
| 436.cactusADM | 165s                 | 194s             | 208s                  |
| 437.leslie3d  | 179s                 | 209s             | 435s                  |
| 454.calculix  | 171s                 | 297s             | 608s                  |
| 459.GemsFDTD  | 261s                 | 286s             | 391s                  |
| 465.tonto     | 295s                 | 373s             | Fails                 |
| 481.wrf       | 157s                 | 271s             | 247s                  |

PGI Fortran: -fast -Mfprelaxed -Mstack\_arrays gfortran: -O3 -funroll-loops -fpeel-loops -ffast-math Flang: -O3 -march=core-avx2 -ffp-contract=fast -Knoieee  
Performance measured November, 2016 and are considered estimates per SPEC run and reporting rules. SPEC® and SPEC CPU® are registered trademarks of the  
Standard Performance Evaluation Corporation ([www.spec.org](http://www.spec.org)).

# Flang OpenMP Performance

SPEC OMP 2012 Fortran codes, all times in seconds, 32 Haswell cores (64 threads)

|            | PGI FORTRAN<br>16.10 | GFORTRAN<br>6.1 | FLANG DEV<br>LLVM 3.9 |
|------------|----------------------|-----------------|-----------------------|
| 350.md     | 517s                 | 3460s           | 459s                  |
| 351.bwaves | 469s                 | 519s            | 805s                  |
| 357.bt331  | 449s                 | 492s            | 474s                  |
| 360.ilbdc  | 541s                 | 6846s           | 539s                  |
| 362.fma3d  | 575s                 | 504s            | 656s                  |
| 363.swim   | 633s                 | 634s            | 632s                  |
| 370.mgrid  | 693s                 | 697s            | 690s                  |
| 371.applu  | 451s                 | 414s            | 514s                  |

PGI Fortran: -fast -mp -Mfprelaxed -Mstack\_arrays gfortran: -O3 -funroll-loops -fpeel-loops -ffast-math -fopenmp

Flang: -O3 -mp -march=core-avx2 -ffp-contract=fast -Knoieee All: OMP\_NUM\_THREADS=64 OMP\_PROC\_BIND=true

Performance measured November, 2016 and are considered estimates per SPEC run and reporting rules. SPEC® and SPEC OMP® are registered trademarks of the Standard Performance Evaluation Corporation ([www.spec.org](http://www.spec.org)).

# Flang Year 2 Development Plans

- Source code
  - Continue source clean-up, refactoring, documentation
  - Create repository and release as open source
  - Deploy an open source testing infrastructure
- Features
  - Enhance compile-time Fortran error/warning messages
  - Incremental F08 and OpenMP 4.5 features
  - LLVM enhancements to enable Fortran DWARF generation
- Performance
  - Incremental, likely to be reactive after initial pass is done

# Concluding Thoughts

- LLVM is integral to HPC compilers at NVIDIA and PGI
- Fortran → First-class citizen in the LLVM community
- LLVM as a platform for out-of-tree developers

